

VA-ATACOM Supplementary Material: The DT-margin Heuristic Baseline

Anonymous submission

Table 1: DT-margin ablation on Goal (5 seeds).

Config	α	h	Mean $\bar{C}$	GO/5
Velocity-adaptive only	0.3	0	2.90	5/5
BRT-only	0.0	3	2.35	4/5
DT-margin (full)	0.3	3	0.96	5/5

The DT-margin Heuristic Baseline

VA-ATACOM’s velocity augmentation has a heuristic predecessor we call the **DT-margin** filter — the “DT-margin (heuristic)” row of Table 1 in the main paper. It inflates the keepout radius linearly with speed, $r_{\text{eff}}(\mathbf{v}) = r_{\text{base}}(1 + \alpha \|\mathbf{v}\|)$, projects out the inward radial component, and adds the same sim-BRT reward shaping. As shown in the main paper’s design study, this linear margin is the *first-order* approximation of VA-ATACOM’s quadratic braking constraint; under fair hazard-radius matching it reaches 13/15 GO — two Push seeds short of VA-ATACOM’s 15/15, illustrating that the linearisation almost captures the right behaviour but leaves a residual gap that only the exact quadratic constraint closes. We characterise it here for completeness — its two tunable knobs (α and the BRT horizon h) and their interaction with the control rate. These findings motivated, and are superseded by, the parameter-free velocity-augmented manifold of the main paper.

Ablation of the two components

Table 1 (Fig. 1) shows both components contribute: velocity-adaptive margin alone reaches 5/5 GO at higher variance (2.90); BRT alone 4/5 GO (one M1 collapse at $\bar{C}=7.98$); the full heuristic 0.96 cost / 100% GO.

Δt sensitivity

Table 2 corroborates the M1 chord-excursion premise of the main paper’s failure-mode analysis (and hence Prop. 4): continuous-time ATACOM achieves zero cost at $\Delta t \leq 0.02\text{s}$ but fails at $\Delta t \geq 0.05\text{s}$, the phase transition matching the M1 threshold $\Delta t \|u\| \gtrsim \sqrt{d_{\text{safe}}} \cdot r$. The DT-margin heuristic tracks the safe regime at $\Delta t \in \{0.10, 0.20\}\text{s}$ with $h=3$. At $\Delta t=0.05\text{s}$ the picture is nuanced: the horizon rule prescribes $h^*=4$, but a 5-seed sweep finds $h=3$ at 3/5 GO (bimodal), $h=4$ at 0/5, and $h=5$ at 3/5 (one seed at 37.05). We attribute

This is an anonymized submission for review purposes only. Distribution, citation, or public sharing of this manuscript is strictly prohibited. Copyright and publication details will appear in the final version if accepted.




Figure 1: DT-margin ablation: both components are necessary. Mean episode cost $\bar{C}$ on SafetyPointGoal1-v0 (5 seeds $\times$ 200K env steps), bars are ± 1 std. VAM-only ($\alpha=0.3, h=0$) addresses M1 chord excursion only; BRT-only ($\alpha=0, h=3$) anticipates M2 drift but one seed collapses via M1; DT-margin ($\alpha=0.3, h=3$) combines both for lowest mean (0.96) and 5/5 GO.

Table 2: DT-margin cost vs. time step ($\bar{C} \pm \sigma$, 5 seeds, 50K steps).

Δt	ATACOM	DT-margin ($h=3$)	Note
0.01s	0.00	3.02	ATACOM OK
0.02s	0.00	0.00	ATACOM OK
0.05s	39.95	4.96 $\pm$ 6.08 (3/5 GO)	Basin regime
0.10s	13.68	2.60	DT-margin OK
0.20s	—	1.33 $\pm$ 1.50	$h \gg h^*$

this to a *BRT-shaping basin*: near the penetration horizon the lookahead reports unavoidable penetration along multiple directions every step, integrating spurious penalty into the policy gradient and trapping training.

α/h sensitivity

Sweeping α and h on Goal (5 seeds $\times$ 200K):

Only configurations at least one step above the minimum on *both* axes ($\alpha \geq 0.1, h \geq 2$) reach 5/5 GO; setting $\alpha=0$ drops to 4/5 with a seed at $\bar{C}=8.46$. The over-provisioned config ($\alpha=0.5, h=5$) *regresses* to 4/5 ($\bar{C}=9.26$ outlier) — large- h BRT shaping can integrate spurious penalty (see the Δt sensitivity above) and destabilise training. The DT-margin recipe

Table 3: DT-margin α/h sensitivity on Goal (5 seeds $\times$ 200K).

Config	α	h	mean $\bar{C}$	GO/5
no margin	0.0	2	3.37 ± 2.81	4/5
short horizon	0.05	1	1.94 ± 1.78	4/5
both at minimum	0.05	2	2.36 ± 2.60	4/5
above minimum	0.1	2	1.30 ± 1.66	5/5
α generous	0.3	1	1.72 ± 1.46	5/5
paper choice	0.3	3	0.96 ± 0.68	5/5
over-provisioned	0.5	5	3.00 ± 3.38	4/5

is thus robust only with margin on both knobs; VA-ATACOM eliminates α entirely, replacing it with the calibrated $a_{\max}$.

Statistical Significance

Safety-Gymnasium (Table 1 of the main paper). We pair every (task, seed) of each baseline against the same (task, seed) of VA-ATACOM (15 paired episode-cost samples per baseline; 11 baselines $\times$ 15 = 165 paired tests in total). One-sided Wilcoxon signed-rank tests (alternative: baseline cost > VA-ATACOM cost) give:

Baseline	p -value	sig.
PPO baseline	3.1×10^{-5}	††
ATACOM (null)	3.1×10^{-5}	††
ATACOM-VD	3.1×10^{-5}	††
ATACOM-S	3.1×10^{-5}	††
ATACOM-LA	3.1×10^{-5}	††
HOCBF	3.1×10^{-5}	††
DCM	3.1×10^{-5}	††
Predictive ATACOM	4.9×10^{-4}	††
CBF-QP	3.1×10^{-5}	††
PPO-Lag	3.1×10^{-5}	††
DT-margin (1st-order approx.)	2.9×10^{-3}	†

†: $p < 0.05$, ††: $p < 0.001$

All 11 baselines are significantly worse than VA-ATACOM ($p < 0.05$); 10 are extremely significant ($p < 0.001$). DT-margin, the first-order Taylor approximation of VA-ATACOM, has the smallest gap ($p=0.003$) — consistent with the main paper’s claim that the linearisation captures most of the safe behaviour but leaves a residual gap that only the exact quadratic constraint closes (10/15 paired cells where DT-margin > VA-ATACOM, 4/15 ties at 0). A Friedman omnibus test across all 12 methods rejects equality with $\chi^2=97.05$, $p=6.8 \times 10^{-16}$ (df=11).

Webots transfer (120 paired trials). For the Webots E-puck study (§ Sim-to-Real of main paper) we apply the McNemar test for paired binary outcomes (deep penetration ≥ 1 cm yes/no) on the same (start, goal, world, controller) configuration with the filter on vs. off:

The filter *saves* 23 paired trials (filter on safe, filter off deep) and *hurts* 0 paired trials (filter on deep, filter off safe). The number-needed-to-filter is $N/\text{saves} = 120/23 \approx 5.2$: on

Setting	N	on-deep	off-deep	saves	p
P-controller (3 worlds)	60	1	14	13	1.2×10^{-4}
PPO transfer (3 worlds)	60	0	10	10	9.8×10^{-4}
All 120 trials	120	1	24	23	1.2×10^{-7}

average one in every five trajectories needs the filter to avoid deep penetration. The one VA-ATACOM deep case (dense, P-controller, triple-obstacle pocket) is *also* deep without the filter, so the filter never makes things worse.

Computing Infrastructure

All experiments reported in the main paper and this supplementary were run on a single workstation:

- **CPU:** Intel Core i7-13850HX (20 cores, 28 threads).
- **RAM:** 31 GiB DDR5.
- **OS:** Ubuntu 24.04 LTS, Linux kernel 6.17.
- **GPU:** not used — all training and evaluation run *single-threaded on CPU* (OMP_NUM_THREADS=1, MKL_NUM_THREADS=1). Each Safety-Gymnasium 200K-step run takes ≈ 15 –20 min wall-clock; we parallelise up to 20 such single-threaded runs at once.

Software stack (versions used at experiment time):

- Python 3.10
- PyTorch 2.x (CPU build)
- NumPy 1.23 / matplotlib 3.x
- safety-gymnasium 1.0 (the Safety-Gym successor; we ran with SafetyPointGoal1-v0, SafetyPointPush1-v0, SafetyPointMultiGoal0-v0)
- gymnasium 0.28 / mujoco 2.3
- Webots R2023b (for the rigid-body validation in main § Sim-to-Real)

Random seeds 0, 1, 2, 3, 4 are used for every (method, task) cell in Table 1 of the main paper. Total compute for the full Table 1 + Webots study is approximately **200 CPU-hours** on the workstation above (~ 10 h wall-clock with 20-way parallelism); the L-shape audit re-run added a further ~ 5 CPU-hours.