

AMRF learns $\alpha(s)$ that is uncorrelated with state — the recipe's fixed $\alpha=0.3$ is in the same plateau

